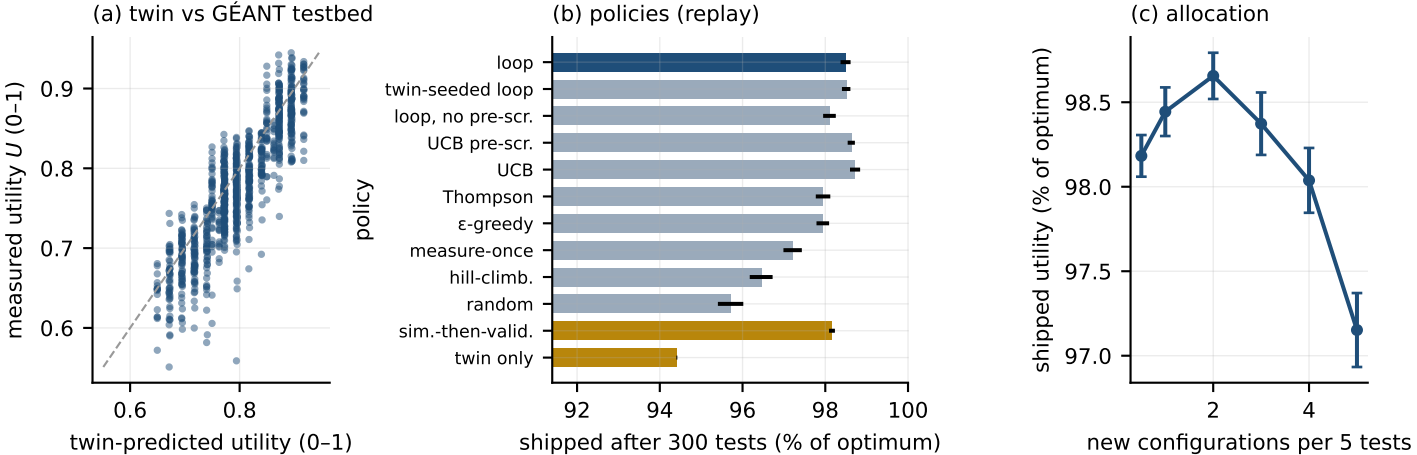
(a) twin vs GÉANT testbed
measured utility U (0–1)
0.9
0.8
0.7
0.6
0.6
0.8
twin-predicted utility (0–1)
(b) policies (replay)
policy
loop
twin-seeded loop
loop, no pre-scr.
UCB pre-scr.
UCB
Thompson
ε-greedy
measure-once
hill-climb.
random
sim.-then-valid.
twin only
92
94
96
98
100
shipped after 300 tests (% of optimum)
(c) allocation
shipped utility (% of optimum)
98.5
98.0
97.5
97.0
2
4
new configurations per 5 tests